


Full data (df=1.0, 100M tokens), d12/w768 base — strategy: init_shuffle (depth scaling)

(A) ensembling reduces loss across training



(B) ensembling beats depth-scaling at matched compute

