


Full data (df=1.0, 100M tokens), d12/w768 base — strategy: init_shuffle (Panel A: per-epoch mean)

(A) ensembling reduces loss across training (per-epoch mean)



(B) ensembling beats width-scaling at matched compute

