

Validation loss vs training, $df = 1.0$, $\lambda = 0$, constant LR, $E = 1$ -- every $\mathcal{L}^*$ above is the minimum of one such curve

