

Validation loss vs training,  $df = 1.0$ ,  $\lambda = 0$ , constant LR,  $E = 1$  -- every  $\mathcal{L}^*$  above is the minimum of one such curve

