



(A) min val loss vs P — fits at various assumed L_0
(init + shuffle shown; init is nearly identical)



(B) r_P identification curve: r_P depends on L_0



(C) overfit-onset position vs P
epoch is constant; tokens scale linearly with P

