

Val loss vs model size (df=1.0, 100M tokens, wd=0)

